



Decision Tree

Nimisha Roy
Georgia Tech



Visual Introduction to Decision Tree



Building a tree to distinguish homes in New York from homes in San Francisco

Decision Tree: Example (2)

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

Outlook:

Sunny,
Overcast,
Rainy

Temperature:

Hot,
Medium,
Cool

Humidity:


High,
Normal,
Low

Wind:

Strong,
Weak

Will I play tennis today?

Decision trees (DT)



The classifier:

$f_T(x)$: majority class in the leaf in the tree T containing x

Model parameters: The tree structure and size

Decision trees

Pieces:

1. Find the **best attribute** to split on
2. Find the **best split** on the **chosen attribute**
3. Decide on when to stop splitting


Categorical or Discrete attributes

- Three variables:
 - Hair = {blond, dark}
 - Height = {tall, short}

Label – Country = {Gromland, Polvia}


Training data:

(B,T,P)
(B,T,P)
(B,S,G)
(D,S,G)
(D,T,G)
(B,S,G)




At each level of the tree,
we split the data
according to the value
of one of the attributes

After enough splits, only
one class is represented
in the node → This is a
terminal leaf of the tree
We call that class the
output class for that
node




General Decision Tree (Discrete Attributes)



Continuous attributes

Decision Tree Example





$X_2 = 0.5$




$X_1 = 0.5$

$X_2 < 0.5??$

$X_1 < 0.5 ??$




Test data



The class of a new input can be classified by following the tree all the way down to a leaf and by reporting the output of the leaf. For example:
 $(0.2, 0.8)$ is classified as $\times$
 $(0.8, 0.2)$ is classified as $\bullet$


General Decision Tree (Continuous Attributes)



Basic Questions


- How to choose the attribute/value to split on at each level of the tree?
- When to stop splitting? When should a node be declared a leaf?
- If a leaf node is impure, how should the class label be assigned?
- If the tree is too large, how can it be pruned?

How to choose the attribute/value to split on at each level of the tree?




- Two classes (red circles/green crosses)
- Two attributes: X_1 and X_2
- 11 points in training data
- Idea $\rightarrow$ Construct a decision tree such that the leaf nodes predict correctly the class for all the training examples

How to choose the attribute/value to split on at each level of the tree?



We want to find the most compact, smallest size tree (Occam's razor), that classifies the training data correctly → We want to find the split choices that will get us the fastest to pure nodes



Information Theory Recap

- $I(x)$
- $H(x)$
- Binary Case: Maximum value of entropy is
- Non-binary case: Maximum value of entropy is
- Range of entropy

Information Content

Coin flip

C_{1H}	0
C_{1T}	6

$$P(C_{1H}) = 0/6 = 0$$

$$P(C_{1T}) = 6/6 = 1$$

C_{2H}	1
C_{2T}	5

$$P(C_{2H}) = 1/6$$

$$P(C_{2T}) = 5/6$$

C_{3H}	2
C_{3T}	4

$$P(C_{3H}) = 2/6$$

$$P(C_{3T}) = 4/6$$

Which coin will give us the purest information?

Entropy ~ Uncertainty

Lower uncertainty, higher information gain

$$H(X) = - \sum_{i=1}^N P(x=i) \log_2 P(x=i)$$

$$\text{Entropy} = - 0 \log 0 - 1 \log 1 = - 0 - 0 = 0$$

$$\text{Entropy} = - (1/6) \log_2 (1/6) - (5/6) \log_2 (5/6) = 0.65$$

$$\text{Entropy} = - (2/6) \log_2 (2/6) - (4/6) \log_2 (4/6) = 0.92$$


Entropy

- In general, the average number of bits necessary to encode n values is the entropy:


$$H = - \sum_{i=1}^n P_i \log_2 P_i$$

- P_i = probability of occurrence of value i
 - High entropy $\rightarrow$ All the classes are (nearly) equally likely
 - Low entropy $\rightarrow$ A few classes are likely; most of the classes are rarely observed

Entropy




The entropy captures the degree of “purity” of the distribution



Example Entropy Calculation

①




$$N_A = 1$$
$$N_B = 6$$

$$p_A = N_A / (N_A + N_B) = 1/7$$
$$p_B = N_B / (N_A + N_B) = 6/7$$

$$H_1 = -p_A \log_2 p_A - p_B \log_2 p_B$$
$$= 0.59$$

②



$$N_A = 3$$
$$N_B = 2$$


$$p_A = N_A / (N_A + N_B) = 3/5$$
$$p_B = N_B / (N_A + N_B) = 2/5$$

$$H_2 = -p_A \log_2 p_A - p_B \log_2 p_B$$
$$= 0.97$$

$$H_1 < H_2 \Rightarrow (2) \text{ less pure than } (1)$$

Conditional Entropy

Entropy before splitting: H



After splitting, a fraction P_L of the data goes to the left node, which has entropy H_L

After splitting, a fraction P_R of the data goes to the right node, which has entropy H_R

The average entropy after splitting is:


Entropy of left node

$$H_L \times P_L + H_R \times P_R$$

“Conditional Entropy”

Probability that a random input is directed to the left node

Information Gain



We want nodes as pure as possible

→ We want to reduce the entropy as much as possible

→ We want to maximize the difference between the entropy of the parent node and the expected entropy of the children

Information Gain (IG) = Amount by which the ambiguity is decreased by splitting the node


Maximize:

$$IG = H - (H_L \times P_L + H_R \times P_R)$$

Notations


- Entropy: $H(Y)$ = Entropy of the distribution of classes at a node
- Conditional Entropy:
 - *Discrete*: $H(Y|X_j)$ = Entropy after splitting with respect to variable j
 - *Continuous*: $H(Y|X_j, t)$ = Entropy after splitting with respect to variable j with threshold t
- Information gain:
 - *Discrete*: $IG(Y|X_j) = H(Y) - H(Y|X_j)$ = Entropy after splitting with respect to variable j
 - *Continuous*: $IG(Y|X_j, t) = H(Y) - H(Y|X_j, t)$ = Entropy after splitting with respect to variable j with threshold t

$$H = 0.99$$




$$IG =$$

$$H - (H_L * 4/11 + H_R * 7/11)$$




$$H_L = 0$$




$$H_R = 0.58$$

$$H = 0.99$$




$$IG =$$


$$H - (H_L * 5/11 + H_R * 6/11)$$




$$H_L = 0.97$$




$$H_R = 0.92$$






More Complete Example



● = 20 training examples from class A


× = 20 training examples from class B

Attributes = X_1 and X_2 coordinates



X_1 Split value


Best split value (max Information Gain) for X_1
attribute: 0.24 with $IG = 0.138$



Best split value (max Information Gain) for X_2 attribute: 0.234 with $IG = 0.202$

Best X_1 split: 0.24, IG = 0.138
Best X_2 split: 0.234, IG = 0.202

Split on X_2 with 0.234



$X_2 \leq 0.233657$

Total data points = 7
7 A
0 B


Total data points = 33
13 A
20 B

Left direction for smaller value, right direction for bigger value

Best X split: 0.24 IG = 0.138

Best Y split: 0.202

There is no point in splitting this node further since it contains only data from a single class $\rightarrow$ return it as a leaf node with output 'A'



This node is not pure so we need to split further

$X_2 \leq 0.233$

Total data points = 7



7 A

0 B

Total data points = 33



13 A

20 B



Best split value (max Information Gain) for X_1 attribute: 0.22 with IG $\sim$ 0.182

IG





X_2 Split value

Best split value (max Information Gain) for X_2 attribute: 0.75 with IG ~ 0.353



Best X_1 split: 0.22, IG = 0.182
Best X_2 split: 0.75, IG = 0.353


Split on X_2 with 0.75




Best X_2 split: 0.0010 0.482
 B 53

There is no point in splitting this node further since it contains only data from a single class $\rightarrow$ return it as a leaf node with output 'A'






Final decision tree




Each of the leaf nodes is pure $\rightarrow$ contains data from only one class



Final decision tree
 Given an input $(X, Y) \rightarrow$
 Follow the tree down to a
 leaf.

Return corresponding
 output class for this leaf

Example $(X, Y) = (0.5, 0.5)$





Basic Questions

- How to choose the attribute/value to split on at each level of the tree?
- • When to stop splitting? When should a node be declared a leaf?
- • If a leaf node is impure, how should the class label be assigned?
- If the tree is too large, how can it be pruned?

When to stop splitting? Common strategies:

1. Pure and impure leave nodes

- All points belong to the same class; OR
- All points from one class completely overlap with points from another class (i.e., same attributes)
 - Output majority class as this leaf's label



2. Node contains points fewer than some threshold

3. Node purity is higher than some threshold

4. Further splits provide no improvement in *training loss*

$$(\text{loss}(T) \leq \text{loss}(T_L) + \text{loss}(T_R))$$

Decision Tree Algorithm (Continuous Attributes)

- LearnTree(X, Y)
 - Input:
 - Set X of R training vectors, each containing the values $(x_1, \dots, x_M)$ of M attributes $(X_1, \dots, X_M)$
 - A vector Y of R elements, where y_j = class of the j^{th} datapoint
 - If all the datapoints in X have the same class value y
 - Return a leaf node that predicts y as output
 - If all the datapoints in X have the same attribute value $(x_1, \dots, x_M)$
 - Return a leaf node that predicts the majority of the class values in Y as output
 - Try all the possible attributes X_j and threshold t and choose the one, j^* , for which $\text{IG}(Y|X_j, t)$ is maximum
 - X_L, Y_L = set of datapoints for which $x_{j^*} < t$ and corresponding classes
 - X_H, Y_H = set of datapoints for which $x_{j^*} \geq t$ and corresponding classes
 - Left Child $\leftarrow \text{LearnTree}(X_L, Y_L)$
 - Right Child $\leftarrow \text{LearnTree}(X_H, Y_H)$

Decision Tree Algorithm (Discrete Attributes)

- LearnTree(X, Y)
 - Input:
 - Set X of R training vectors, each containing the values $(x_1, \dots, x_M)$ of M attributes $(X_1, \dots, X_M)$
 - A vector Y of R elements, where y_j = class of the j^{th} datapoint
 - If all the datapoints in X have the same class value y
 - Return a leaf node that predicts y as output
 - If all the datapoints in X have the same attribute value $(x_1, \dots, x_M)$
 - Return a leaf node that predicts the majority of the class values in Y as output
 - Try all the possible attributes X_j and choose the one, j^* , for which $IG(Y|X_j)$ is maximum
 - For every possible value v of X_{j^*} :
 - X_v, Y_v = set of datapoints for which $x_{j^*} = v$ and corresponding classes
 - $\text{Child}_v \leftarrow \text{LearnTree}(X_v, Y_v)$

Decision Trees So Far

- Given N observations from training data, each with D attributes X and a class attribute Y , construct a sequence of tests (decision tree) to predict the class attribute Y from the attributes X
- Basic strategy for defining the tests (“when to split”) → maximize the information gain on the training data set at each node of the tree
- Problems (next):
 - Computational issues → How expensive is it to compute the IG
 - The tree will end up being much too big → *pruning*
 - Evaluating the tree on training data is dangerous → *overfitting*

Basic Questions

- How to choose the attribute/value to split on at each level of the tree?
- When to stop splitting? When should a node be declared a leaf?
- If a leaf node is impure, how should the class label be assigned?
- ➔ • If the tree is too large, how can it be pruned?

What will happen if a tree is too large?

Overfitting

High variance


Instability in predicting test data

How to avoid overfitting?

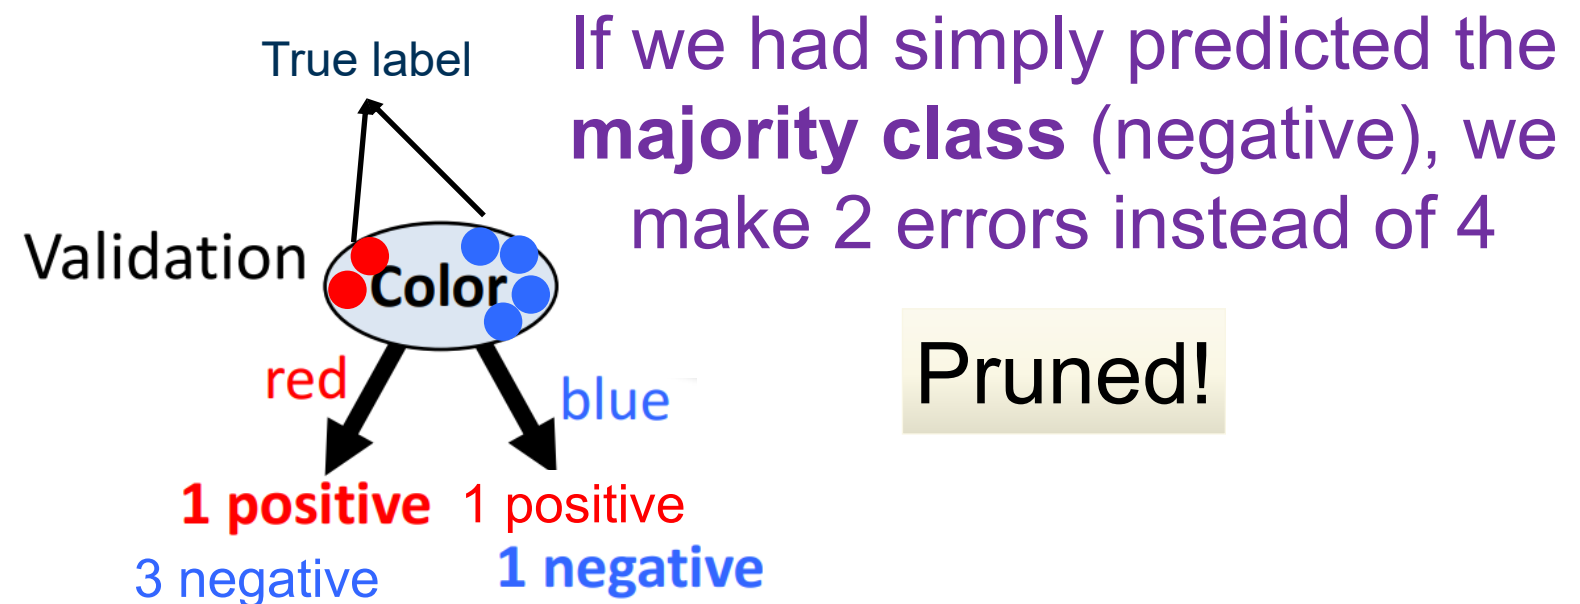
- Acquire more training data
- Remove irrelevant attributes (manual process – not always possible)
- Grow full tree, then post-prune
- Ensemble learning

How to decide to remove it a node using pruning

- Pruning of the decision tree is done by replacing a whole subtree by a leaf node.
- The replacement takes place if a decision rule establishes that the expected error rate in the subtree is greater than in the single leaf.



3 training data points
Actual Label: 1 positive and 2 negative
Predicted Label: 1 positive and 2 negative
3 correct and 0 incorrect



If we had simply predicted the **majority class** (negative), we make 2 errors instead of 4

Pruned!

6 validation data points
Actual label: 2 positive and 4 negative
Predicted Label: 4 positive and 2 negative
2 correct and 4 incorrect

Reduced-Error Pruning

Split data into training and validation sets

Grow tree based on training set

Do until further pruning is harmful:

1. Evaluate impact on validation set of pruning each possible node (plus those below it)
2. Greedily remove the node that most improves validation set accuracy